

EXPT – 2 : Email Spam or Ham Classification

REPORT BY GOPIKA GANESAN

AIM:

Email Spam or Ham Classification using Naive Bayes, KNN, and SVM.

LIBRARIES USED:

NumPy, pandas, scikit learn, seaborn, matplotlib

OBJECTIVE:

To classify emails as spam or ham using three classification algorithms—Naive Bayes, K-Nearest Neighbors (KNN), and Support Vector Machine (SVM)—and evaluate their performance using accuracy metrics and K-Fold cross-validation.

CODE FOR ALL VARIANTS AND MODELS:

1.Naive Bayes

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer

from sklearn.metrics import (
    confusion_matrix, classification_report, roc_auc_score, roc_curve,
    accuracy_score, precision_score, recall_score, f1_score, fbeta_score
)

from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB

# 3. EDA & Handling Missing Values
print("\nMissing Values:\n")
print(df.isnull().sum())
target="class"
# 4. Outlier Detection Function
def detect_outliers(df, col):

```

```

def detect_outliers(df, col):
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    return df[(df[col] < Q1 - 1.5 * IQR) | (df[col] > Q3 + 1.5 * IQR)]

# 5. Train-Test Split
X = df.drop(columns=target)
y = df[target]

X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size=0.25, random_state=42)

# 6. Train and Evaluate
models = {
    "GaussianNB": GaussianNB(),
    "MultinomialNB": MultinomialNB(),
    "BernoulliNB": BernoulliNB()
}

gnb = GaussianNB()
gnb.fit(X_train, y_train)
y_pred_gnb = gnb.predict(X_val)

```

```

print("\n GaussianNB ")
print("Accuracy:", accuracy_score(y_val, y_pred_gnb))
print("Precision:", precision_score(y_val, y_pred_gnb, average='weighted'))
print("Recall:", recall_score(y_val, y_pred_gnb, average='weighted'))
print("F1 Score:", f1_score(y_val, y_pred_gnb, average='weighted'))
print("F-beta Score ( $\beta=0.5$ ):", fbeta_score(y_val, y_pred_gnb, beta=0.5, average='weighted'))

print("\nConfusion Matrix:\n", confusion_matrix(y_val, y_pred_gnb))
print("\nClassification Report:\n", classification_report(y_val, y_pred_gnb))

if hasattr(gnb, "predict_proba"):
    y_proba_gnb = gnb.predict_proba(X_val)[:, 1]
    print("ROC AUC Score:", roc_auc_score(y_val, y_proba_gnb))
    fpr, tpr, _ = roc_curve(y_val, y_proba_gnb)
    plt.plot(fpr, tpr, label="GaussianNB")

mnb = MultinomialNB()
mnb.fit(X_train, y_train)
y_pred_mnb = mnb.predict(X_val)

print("\n MultinomialNB ")
print("Accuracy:", accuracy_score(y_val, y_pred_mnb))
print("Precision:", precision_score(y_val, y_pred_mnb, average='weighted'))

```

```

print("Accuracy:", accuracy_score(y_val, y_pred_mnb))
print("Precision:", precision_score(y_val, y_pred_mnb, average='weighted'))
print("Recall:", recall_score(y_val, y_pred_mnb, average='weighted'))
print("F1 Score:", f1_score(y_val, y_pred_mnb, average='weighted'))
print("F-beta Score ( $\beta=0.5$ ):", fbeta_score(y_val, y_pred_mnb, beta=0.5, average='weighted'))

print("\nConfusion Matrix:\n", confusion_matrix(y_val, y_pred_mnb))
print("\nClassification Report:\n", classification_report(y_val, y_pred_mnb))

if hasattr(mnb, "predict_proba"):
    y_proba_mnb = mnb.predict_proba(X_val)[:, 1]
    print("ROC AUC Score:", roc_auc_score(y_val, y_proba_mnb))
    fpr, tpr, _ = roc_curve(y_val, y_proba_mnb)
    plt.plot(fpr, tpr, label="MultinomialNB")

bnb = BernoulliNB()
bnb.fit(X_train, y_train)
y_pred_bnb = bnb.predict(X_val)

print("\n BernoulliNB ")
print("Accuracy:", accuracy_score(y_val, y_pred_bnb))
print("Precision:", precision_score(y_val, y_pred_bnb, average='weighted'))
print("Recall:", recall_score(y_val, y_pred_bnb, average='weighted'))

```

```

print("F1 Score:", f1_score(y_val, y_pred_bnb, average='weighted'))
print("F-beta Score ( $\beta=0.5$ ):", fbeta_score(y_val, y_pred_bnb, beta=0.5, average='weighted'))

print("\nConfusion Matrix:\n", confusion_matrix(y_val, y_pred_bnb))
print("\nClassification Report:\n", classification_report(y_val, y_pred_bnb))

if hasattr(bnb, "predict_proba"):
    y_proba_bnb = bnb.predict_proba(X_val)[:, 1]
    print("ROC AUC Score:", roc_auc_score(y_val, y_proba_bnb))
    fpr, tpr, _ = roc_curve(y_val, y_proba_bnb)
    plt.plot(fpr, tpr, label="BernoulliNB")

plt.title("ROC Curves")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()

```

2.KNN:

```

X = df.drop(columns="class")
y = df["class"]

X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size=0.25, random_state=42)

print("\n KDTTree Algorithm (k=5)")
knn_kdtree = Pipeline([
    ("scaler", StandardScaler()),
    ("knn", KNeighborsClassifier(n_neighbors=5, algorithm="kd_tree"))
])
start_time = time.time()
knn_kdtree.fit(X_train, y_train)
training_time=time.time()-start_time
y_pred = knn_kdtree.predict(X_val)

print(f"Training Time: {training_time:.4f} seconds")
print("Accuracy:", accuracy_score(y_val, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))

```

```

print("\n BallTree Algorithm (k=5)")
knn_balltree = Pipeline([
    ("scaler", StandardScaler()),
    ("knn", KNeighborsClassifier(n_neighbors=5, algorithm="ball_tree"))
])
start_time = time.time()
knn_balltree.fit(X_train, y_train)
training_time = time.time() - start_time
y_pred = knn_balltree.predict(X_val)

print(f"Training Time: {training_time:.4f} seconds")
print("Accuracy:", accuracy_score(y_val, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))

```

```

print("\n (Varying K)")
for k in [3, 5, 7, 9]:
    knn_auto = Pipeline([
        ("scaler", StandardScaler()),
        ("knn", KNeighborsClassifier(n_neighbors=k, algorithm="auto"))
    ])
    knn_auto.fit(X_train, y_train)

```

```

y_pred = knn_auto.predict(X_val)

```

```

print(f"\n K = {k} ")
print("Accuracy:", accuracy_score(y_val, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))

```

3.SVM

```

X = df.drop(columns="class")
y = df["class"]
classes = y.unique()
y_bin = label_binarize(y, classes=classes)
n_classes = y_bin.shape[1] if y_bin.ndim > 1 else 1
X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size=0.25, random_state=42)

# Kernels to try
kernels = ['linear', 'poly', 'rbf', 'sigmoid']

for kernel in kernels:
    print(f"\nSVM with {kernel} kernel")
    svm_pipeline = Pipeline([
        ("scaler", StandardScaler()),
        ("svm", SVC(kernel=kernel, probability=True)) # probability=True needed for ROC
    ])

    start_time = time.time()
    svm_pipeline.fit(X_train, y_train)
    training_time = time.time() - start_time

    y_pred = svm_pipeline.predict(X_val)
    y_proba = svm_pipeline.predict_proba(X_val)

```

```

print(f"Training time: {training_time:.4f} seconds ")
print("Accuracy:", accuracy_score(y_val, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))
y_val_bin = label_binarize(y_val, classes=classes)
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_val_bin[:, i], y_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

plt.figure()
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f'Class {classes[i]} (AUC = {roc_auc[i]:.2f})')
plt.plot([0, 1], [0, 1], 'k--', lw=1)
plt.title(f"ROC Curve - SVM ({kernel} kernel)")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc="lower right")
plt.grid()
plt.show()

```

```

kfold = KFold(n_splits=5, shuffle=True, random_state=42)
svm_cv_pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("svm", SVC(kernel='linear'))
])
cv_scores = cross_val_score(svm_cv_pipeline, X, y, cv=kfold, scoring='accuracy')
print("Cross-Validation Accuracies:", cv_scores)
print("Mean CV Accuracy: {:.4f}".format(np.mean(cv_scores)))

```

```

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import loguniform, randint

# Base pipeline
pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("svm", SVC(probability=True))
])

# Parameter spaces for each kernel
param_spaces = {
    "linear": {"svm__kernel": ["linear"], "svm__C": loguniform(1e-3, 1e3)},
    "poly": {"svm__kernel": ["poly"], "svm__C": loguniform(1e-3, 1e3),
            "svm__degree": randint(2, 6), "svm__gamma": loguniform(1e-4, 1e1)},
    "rbf": {"svm__kernel": ["rbf"], "svm__C": loguniform(1e-3, 1e3),
           "svm__gamma": loguniform(1e-4, 1e1)},
    "sigmoid": {"svm__kernel": ["sigmoid"], "svm__C": loguniform(1e-3, 1e3),
               "svm__gamma": loguniform(1e-4, 1e1)}
}

best_params = {}
for kernel, params in param_spaces.items():
    print(f"\n=== {kernel.upper()} Kernel ===")
    search = RandomizedSearchCV(
        pipeline, param_distributions=params,
        n_iter=10, cv=3, scoring='accuracy',
        n_jobs=-1, random_state=42

```

```

        n_jobs=-1, random_state=42
    )
    search.fit(X_train, y_train)
    best_params[kernel] = search.best_params_
    print("Best Params:", search.best_params_)
    print("Best CV Accuracy:", search.best_score_)

# Validation
y_pred = search.best_estimator_.predict(X_val)
print("Validation Accuracy:", accuracy_score(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))

print("\nSummary of Best Parameters:", best_params)

```

```
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, classification_report
import numpy as np
```

```
# Pipeline
```

```
pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("svm", SVC(probability=True))
])
```

```
# Grid parameters
```

```
param_grid = [
    {"svm__kernel": ["linear"], "svm__C": np.logspace(-3, 3, 5)},
    {"svm__kernel": ["poly"], "svm__C": np.logspace(-3, 3, 5),
     "svm__degree": [2, 3, 4], "svm__gamma": ["scale", "auto"]},
    {"svm__kernel": ["rbf"], "svm__C": np.logspace(-3, 3, 5),
     "svm__gamma": ["scale", "auto"]},
    {"svm__kernel": ["sigmoid"], "svm__C": np.logspace(-3, 3, 5),
     "svm__gamma": ["scale", "auto"]}
]
```

```
# Grid search
```

```
grid = GridSearchCV(pipeline, param_grid, cv=3, scoring="accuracy", n_jobs=-1)
grid.fit(X_train, y_train)
```

```
print("Best Parameters:", grid.best_params_)
print("Best CV Accuracy:", grid.best_score_)
```

```
# Validation
```

```
y_pred = grid.best_estimator_.predict(X_val)
print("Validation Accuracy:", accuracy_score(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))
```

CONFUSION MATRIX AND ROC FOR EACH:

1.Naïve Bayes:

GaussianNB

Accuracy: 0.8358695652173913

✓ Precision: 0.8666295251955612

Recall: 0.8358695652173913

F1 Score: 0.8377216632152633

F-beta Score ($\beta=0.5$): 0.8517900866226876

Confusion Matrix:

[[425 134]

[17 344]]

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.76	0.85	559
1	0.72	0.95	0.82	361
accuracy			0.84	920
macro avg	0.84	0.86	0.83	920
weighted avg	0.87	0.84	0.84	920

ROC AUC Score: 0.9545215784022716

MultinomialNB

Accuracy: 0.8

Precision: 0.7984079601990048

Recall: 0.8

F1 Score: 0.7985370950888192

F-beta Score ($\beta=0.5$): 0.7982959114000469

Confusion Matrix:

[[480 79]

[105 256]]

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.86	0.84	559
1	0.76	0.71	0.74	361
accuracy			0.80	920
macro avg	0.79	0.78	0.79	920
weighted avg	0.80	0.80	0.80	920

ROC AUC Score: 0.8662803086239277

BernoulliNB

Accuracy: 0.8771739130434782

Precision: 0.8766807714723064

Recall: 0.8771739130434782

F1 Score: 0.8766128740859677

F-beta Score ($\beta=0.5$): 0.8765770388407951


Confusion Matrix:

```
[[511  48]
 [ 65 296]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.91	0.90	559
1	0.86	0.82	0.84	361
accuracy			0.88	920
macro avg	0.87	0.87	0.87	920
weighted avg	0.88	0.88	0.88	920

ROC AUC Score: 0.9488426602708634



2.KNN

KDTree Algorithm (k=5)

Training Time: 0.0596 seconds

Accuracy: 0.908695652173913


Confusion Matrix:

```
[[528  31]
```

```
 [ 53 308]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.94	0.93	559
1	0.91	0.85	0.88	361
accuracy			0.91	920
macro avg	0.91	0.90	0.90	920
weighted avg	0.91	0.91	0.91	920



BallTree Algorithm (k=5)

Training Time: 0.0728 seconds

Accuracy: 0.908695652173913

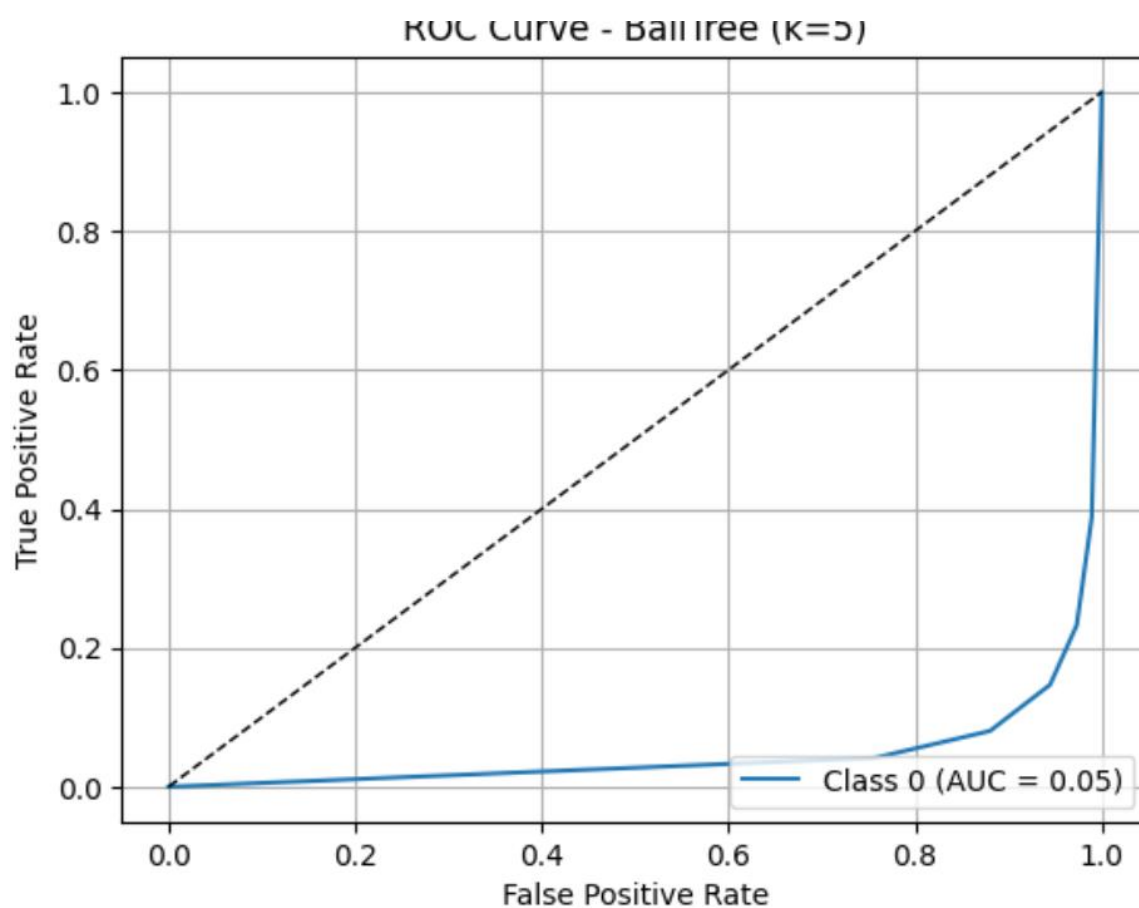
Confusion Matrix:

[[528 31]

[53 308]]

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.94	0.93	559
1	0.91	0.85	0.88	361
accuracy			0.91	920
macro avg	0.91	0.90	0.90	920
weighted avg	0.91	0.91	0.91	920



K = 3

Accuracy: 0.9119565217391304

Confusion Matrix:

[[524 35]

[46 315]]

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.94	0.93	559
1	0.90	0.87	0.89	361
accuracy			0.91	920
macro avg	0.91	0.90	0.91	920
weighted avg	0.91	0.91	0.91	920

K = 5

Accuracy: 0.908695652173913

Confusion Matrix:

[[528 31]

[53 308]]

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.94	0.93	559
1	0.91	0.85	0.88	361
accuracy			0.91	920
macro avg	0.91	0.90	0.90	920
weighted avg	0.91	0.91	0.91	920

K = 7

Accuracy: 0.9021739130434783

Confusion Matrix:

[[526 33]

[57 304]]

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.94	0.92	559
1	0.90	0.84	0.87	361
accuracy			0.90	920
macro avg	0.90	0.89	0.90	920
weighted avg	0.90	0.90	0.90	920

3.SVC



SVM with linear kernel

Training Time: 1.6313 seconds

Accuracy: 0.925

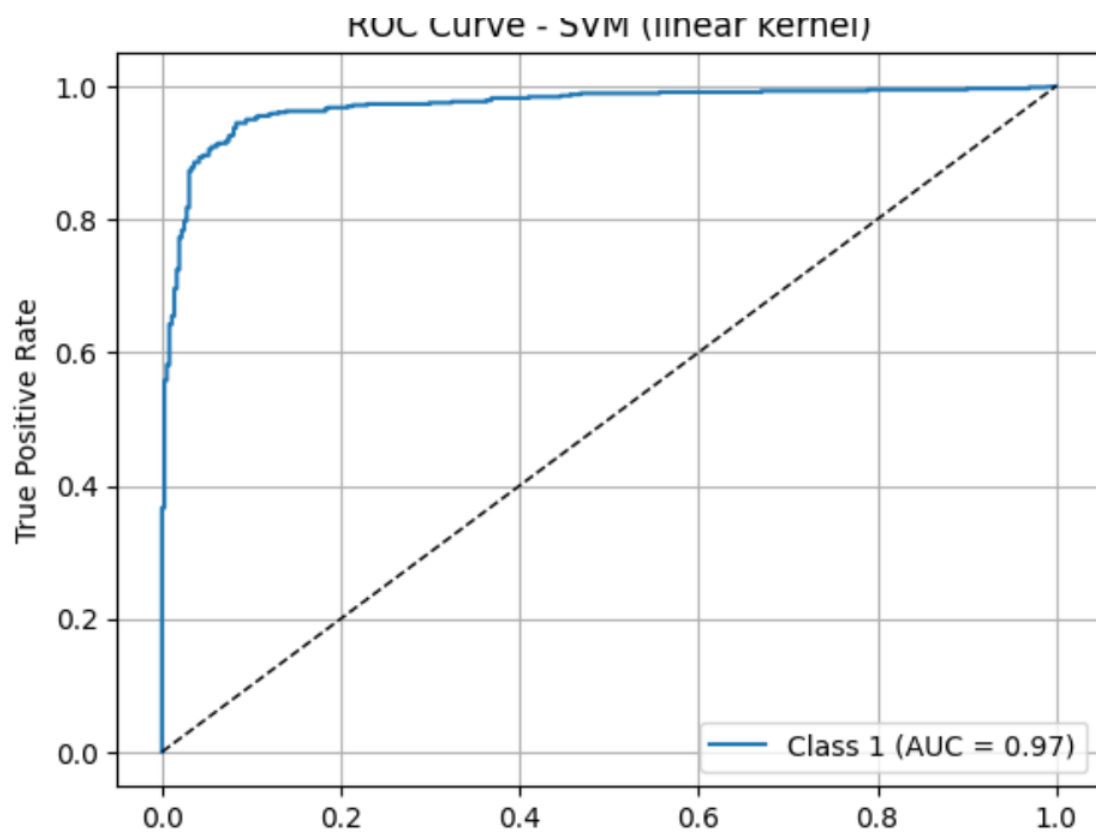
Confusion Matrix:

```
[[536 23]
```

```
[ 46 315]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.96	0.94	559
1	0.93	0.87	0.90	361
accuracy			0.93	920
macro avg	0.93	0.92	0.92	920
weighted avg	0.93	0.93	0.92	920





```
SVM with poly kernel
Training Time: 1.8258 seconds
Accuracy: 0.7760869565217391
Confusion Matrix:
```


```
[[553  6]
 [200 161]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.73	0.99	0.84	559
1	0.96	0.45	0.61	361
accuracy			0.78	920
macro avg	0.85	0.72	0.73	920
weighted avg	0.82	0.78	0.75	920



ROC Curve - SVM (poly kernel)



SVM with rbf kernel

Training Time: 1.1693 seconds

Accuracy: 0.9271739130434783


Confusion Matrix:

```
[[538  21]
```

```
 [ 46 315]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.96	0.94	559
1	0.94	0.87	0.90	361
accuracy			0.93	920
macro avg	0.93	0.92	0.92	920
weighted avg	0.93	0.93	0.93	920



is  Terminal



SVM with sigmoid kernel

Training Time: 1.3583 seconds

✓ Accuracy: 0.8847826086956522

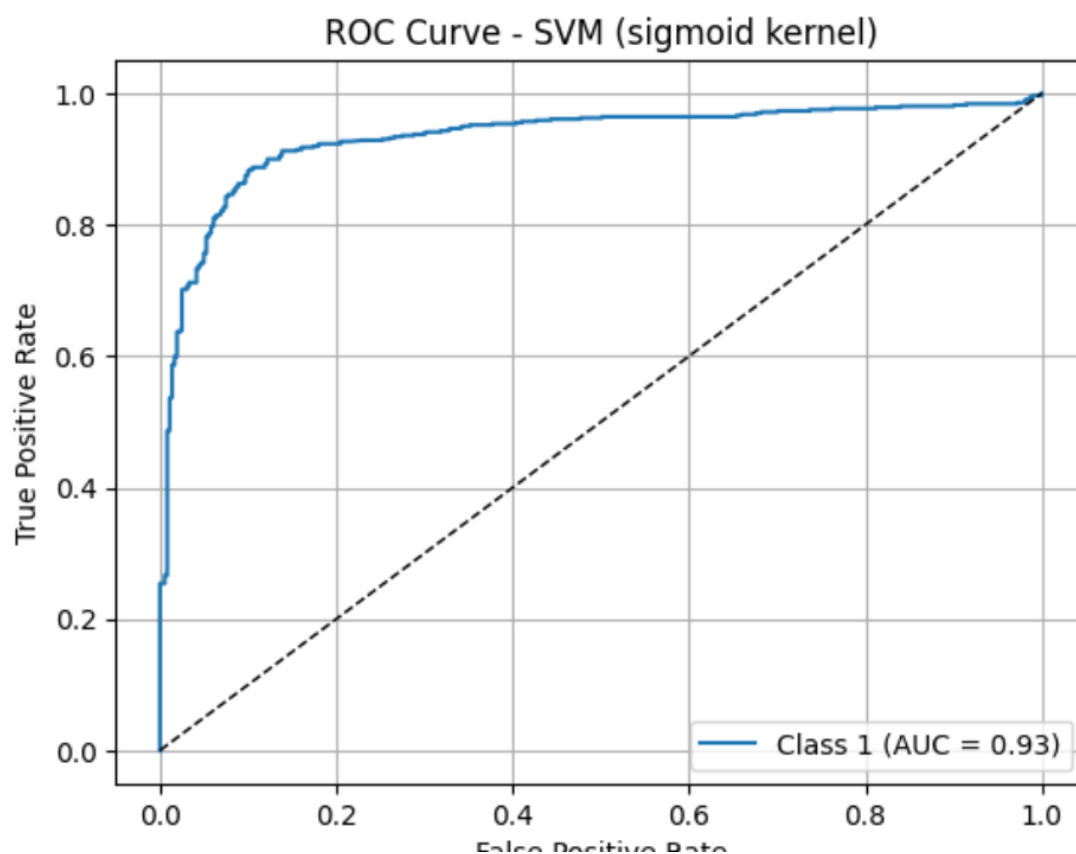
Confusion Matrix:

[[511 48]

[58 303]]

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.91	0.91	559
1	0.86	0.84	0.85	361
accuracy			0.88	920
macro avg	0.88	0.88	0.88	920
weighted avg	0.88	0.88	0.88	920



K-Fold Cross Validation (k=5) using Linear Kernel

Cross-Validation Accuracies: [0.92616721 0.92717391 0.91630435 0.93804348 0.93043478]

Mean CV Accuracy: 0.9276


```
=== LINEAR Kernel ===
Best Params: {'svm__C': np.float64(157.41890047456639), 'svm__kernel': 'linear'}
Best CV Accuracy: 0.9301598570854347
Validation Accuracy: 0.9239130434782609
Classification Report:
```

	precision	recall	f1-score	support
0	0.92	0.96	0.94	559
1	0.94	0.86	0.90	361
accuracy			0.92	920
macro avg	0.93	0.91	0.92	920
weighted avg	0.92	0.92	0.92	920

```
=== POLY Kernel ===
Best Params: {'svm__C': np.float64(0.10051981180656774), 'svm__degree': 5, 'svm__gamma': np.float64(0.3470266988650412), 'svm__kernel': 'poly'}
Best CV Accuracy: 0.9018962504426195
Validation Accuracy: 0.9760869565217392
Classification Report:
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	559
1	0.98	0.96	0.97	361
accuracy			0.98	920
macro avg	0.98	0.97	0.97	920
weighted avg	0.98	0.98	0.98	920

```
=== RBF Kernel ===
Best Params: {'svm__C': np.float64(98.77700294007911), 'svm__gamma': np.float64(0.0011526449540315614), 'svm__kernel': 'rbf'}
Best CV Accuracy: 0.9372264789029509
Validation Accuracy: 0.9456521739130435
Classification Report:
```

	precision	recall	f1-score	support
0	0.94	0.97	0.96	559
1	0.95	0.91	0.93	361
accuracy			0.95	920
macro avg	0.95	0.94	0.94	920
weighted avg	0.95	0.95	0.95	920

```
=== SIGMOID Kernel ===
Best Params: {'svm__C': np.float64(98.77700294007911), 'svm__gamma': np.float64(0.0011526449540315614), 'svm__kernel': 'sigmoid'}
Best CV Accuracy: 0.9149428328442916
Validation Accuracy: 0.9184782608695652
Classification Report:
```

	precision	recall	f1-score	support
0	0.91	0.96	0.93	559
1	0.93	0.86	0.89	361
accuracy			0.92	920
macro avg	0.92	0.91	0.91	920
weighted avg	0.92	0.92	0.92	920

```
Best Parameters: {'svm__C': np.float64(1.0), 'svm__gamma': 'scale', 'svm__kernel': 'rbf'}
Best CV Accuracy: 0.9290740821989202
Validation Accuracy: 0.9434782608695652
Classification Report:
```

	precision	recall	f1-score	support
0	0.94	0.97	0.95	559
1	0.95	0.90	0.93	361
accuracy			0.94	920
macro avg	0.94	0.94	0.94	920
weighted avg	0.94	0.94	0.94	920

COMPARISON TABLES:

Naïve Bayes Variant Comparison:

Metric	Gaussian NB	Multinomial NB	Bernoulli NB
Accuracy	0.836	0.8	0.877
Precision	0.866	0.799	0.877
Recall	0.836	0.8	0.877
F1 Score	0.838	0.799	0.877

KNN: Varying k values:

K	Accuracy	Precision	Recall	F1 Score
1	0.921	0.92	0.92	0.92
3	0.912	0.91	0.90	0.91
5	0.901	0.91	0.90	0.90
7	0.902	0.90	0.89	0.90

KNN: KDTree vs BallTree:

Metric	KDTree	BallTree
Accuracy	0.909	0.909
Precision	0.91	0.91
Recall	0.90	0.90
F1 Score	0.90	0.90
Training Time (s)	0.0357	0.0244

K-Fold Cross-Validation Results (K = 5):

Fold	Naive Bayes Accuracy	KNN Accuracy	SVM Accuracy
Fold 1	0.820	0.893	0.926

Fold 2	0.817	0.908	0.927
Fold 3	0.801	0.925	0.916
Fold 4	0.820	0.906	0.938
Fold 5	0.835	0.908	0.930
Average	0.819	0.908	0.927

SVM Performance with Different Kernels and Parameters

Kernel	Hyperparameters	Accuracy	F1 Score	Training Time
Linear	C=157.419	0.930	0.92	
Polynomial	C=0.101, degree=5, gamma=0.347	0.902	0.97	
RBF	C=98.77, gamma=0.0011	0.94	0.94	
Sigmoid	C=98.77, gamma=0.001,	0.92	0.91	

OBSERVATIONS:

1. Which classifier had the best average accuracy?

SVM performed best with an average accuracy of 0.927, compared to KNN (0.908) and Naïve Bayes (0.819).

2. Which Naïve Bayes variant worked best?

Bernoulli NB was the best, achieving 0.877 accuracy, higher than Gaussian NB (0.836) and Multinomial NB (0.800).

3. How did KNN accuracy vary with k and tree type?

Accuracy decreased slightly as k increased: k=1 → 0.921

- k=3 → 0.912
- k=5 → 0.901

- $k=7 \rightarrow 0.902$

KDTree vs BallTree gave identical accuracy (0.909), but BallTree trained faster (0.024s vs 0.036s).

4. Which SVM kernel was most effective?

The RBF kernel had the best performance overall with 0.94 accuracy and balanced precision/recall.

Linear kernel: 0.930 accuracy

Polynomial kernel: 0.902 accuracy (higher F1 of 0.97 but lower accuracy)

Sigmoid kernel: 0.92 accuracy.

5. How did hyperparameters influence performance?

- For SVM, tuning C and gamma had a strong effect:
 - High C ($\approx 98-157$) improved Linear and RBF accuracy.
 - Polynomial kernel with degree=5 underperformed in accuracy despite good F1.
- For KNN, smaller k gave higher accuracy, while higher k smoothed predictions but reduced accuracy.
- For Naïve Bayes, model choice (distribution assumption) itself acted as the key hyperparameter: Bernoulli worked best for this dataset