

EXPT 4: Ensemble Prediction & Decision Tree Model Evaluation

REPORT BY GOPIKA GANESAN

AIM:

Ensemble Prediction and Decision Tree Model Evaluation for the Wisconsin dataset.

LIBRARIES USED:

NumPy, pandas, scikit learn, seaborn, matplotlib

OBJECTIVE:

To build classifiers such as Decision Tree, AdaBoost, Gradient Boosting, XGBoost, Random Forest, and Stacked Models (using SVM, Naïve Bayes, Decision Tree) and evaluate their performance through 5-Fold Cross-Validation and hyperparameter tuning.

CODE FOR ALL VARIANTS AND MODELS:

1.Decision Tree

```
# Load dataset
data = load_breast_cancer()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target

# Preprocessing
X = StandardScaler().fit_transform(df.drop('target', axis=1))
y = df['target']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Hyperparameter tuning
params = {
    "criterion": ["gini", "entropy"],
    "max_depth": [3, 5, 7, None],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4]
}
grid = GridSearchCV(DecisionTreeClassifier(random_state=42), params, cv=5, scoring="accuracy", n_jobs=-1)
grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("Best Parameters:", grid.best_params_)

# 5-Fold Cross Validation
```

```

fold = 1
cv_scores = []

for train_idx, test_idx in kf.split(X):
    best_model.fit(X[train_idx], y.iloc[train_idx])
    preds = best_model.predict(X[test_idx])
    acc = accuracy_score(y.iloc[test_idx], preds)
    print(f"Fold {fold} Accuracy: {acc:.4f}")
    cv_scores.append(acc)
    fold += 1

print(f"Average 5-Fold CV Accuracy: {np.mean(cv_scores):.4f}")

# Final Test Set Evaluation
y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

print("\n=== Performance Metrics on Test Set ===")
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_prob))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_prob)

```

```

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_prob)
plt.figure(figsize=(6, 4))
plt.plot(fpr, tpr, label=f"AUC = {roc_auc_score(y_test, y_prob):.2f}")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Decision Tree ROC Curve")
plt.legend()
plt.show()

```

2.Adaboost:

```
# 1. Load Dataset
data = load_breast_cancer()
X = data.data
y = data.target
# Split train-test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# 2. Hyperparameter Tuning with GridSearchCV
params = {
    "n_estimators": [50, 100, 150],
    "learning_rate": [0.01, 0.1, 1]
}

grid = GridSearchCV(
    AdaBoostClassifier(random_state=42),
    params,
    cv=5,
    scoring="accuracy",
    n_jobs=-1
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("Best Parameters:", grid.best_params_)
```

```
# 3. 5-Fold Cross-Validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)
fold_accuracies = []
fold = 1

for train_index, val_index in kf.split(X_train):
    X_tr, X_val = X_train[train_index], X_train[val_index]
    y_tr, y_val = y_train[train_index], y_train[val_index]

    best_model.fit(X_tr, y_tr)
    y_pred = best_model.predict(X_val)
    acc = accuracy_score(y_val, y_pred)
    print(f"Fold {fold} Accuracy: {acc:.4f}")
    fold_accuracies.append(acc)
    fold += 1

print("\nAverage CV Accuracy:", np.mean(fold_accuracies))

# 4. Final Evaluation on Test Set
y_pred_test = best_model.predict(X_test)
print("\n=== Performance Metrics on Test Set ===")
print("Accuracy :", accuracy_score(y_test, y_pred_test))
print("Precision:", precision_score(y_test, y_pred_test))
print("Recall   :", recall_score(y_test, y_pred_test))
print("F1 Score :", f1_score(y_test, y_pred_test))
print("ROC-AUC  :", roc_auc_score(y_test, best_model.predict_proba(X_test)[:,:1]))
print("\nClassification Report:\n", classification_report(y_test, y_pred_test))
```

3.Random Forest

```
# 1. Load Dataset
data = load_breast_cancer()
X = data.data
y = data.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# 2. Hyperparameter Tuning
params = {
    "n_estimators": [50, 100, 200],
    "max_depth": [3, 5, 7, None],
    "min_samples_split": [2, 5, 10]
}

grid = GridSearchCV(
    RandomForestClassifier(random_state=42),
    params,
    cv=5,
    scoring="accuracy",
    n_jobs=-1
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("Best Parameters:", grid.best_params_)

# 3. 5-Fold Cross-Validation + ROC
kf = KFold(n_splits=5, shuffle=True, random_state=42)

fold accuracies = []
plt.figure(figsize=(8, 6))

for fold, (train_index, val_index) in enumerate(kf.split(X_train), start=1):
    X_tr, X_val = X_train[train_index], X_train[val_index]
    y_tr, y_val = y_train[train_index], y_train[val_index]

    best_model.fit(X_tr, y_tr)
    y_pred = best_model.predict(X_val)
    y_proba = best_model.predict_proba(X_val)[:, 1]
    acc = accuracy_score(y_val, y_pred)
    fold_accuracies.append(acc)
    print(f"Fold {fold} Accuracy: {acc:.4f}")

    fpr, tpr, _ = roc_curve(y_val, y_proba)
    auc = roc_auc_score(y_val, y_proba)
    plt.plot(fpr, tpr, label=f"Fold {fold} (AUC={auc:.2f})")

print("\nAverage CV Accuracy:", np.mean(fold_accuracies))

# 4. Final Test Evaluation
y_pred_test = best_model.predict(X_test)
y_proba_test = best_model.predict_proba(X_test)[:, 1]

print("\n=== Performance Metrics on Test Set ===")
print("Accuracy :", accuracy_score(y_test, y_pred_test))
print("Precision:", precision_score(y_test, y_pred_test))
print("Recall    :", recall_score(y_test, y_pred_test))
print("F1 Score  :", f1_score(y_test, y_pred_test))
```

```

# -----
# 5. Final ROC Curve for Test Set
# -----
fpr_test, tpr_test, _ = roc_curve(y_test, y_proba_test)
auc_test = roc_auc_score(y_test, y_proba_test)

plt.plot(fpr_test, tpr_test, color="black", linewidth=2.5, label=f"Test ROC (AUC={auc_test:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Random Forest ROC Curve (5-Fold + Test)")
plt.legend()
plt.show()

```

4.GradientBoosting

```

# 1. Load Dataset
data = load_breast_cancer()
X = data.data
y = data.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# 2. Hyperparameter Tuning
params = {
    "n_estimators": [50, 100, 150],
    "learning_rate": [0.01, 0.1, 0.2],
    "max_depth": [3, 5, 7]
}

grid = GridSearchCV(
    GradientBoostingClassifier(random_state=42),
    params,
    cv=5,
    scoring="accuracy",
    n_jobs=-1
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("Best Parameters:", grid.best_params_)
# 3. 5-Fold Cross-Validation + ROC

```

```

kf = KFold(n_splits=5, shuffle=True, random_state=42)
fold accuracies = []
plt.figure(figsize=(8, 6))

for fold, (train_index, val_index) in enumerate(kf.split(X_train), start=1):
    X_tr, X_val = X_train[train_index], X_train[val_index]
    y_tr, y_val = y_train[train_index], y_train[val_index]

    best_model.fit(X_tr, y_tr)
    y_pred = best_model.predict(X_val)
    y_proba = best_model.predict_proba(X_val)[:, 1]
    acc = accuracy_score(y_val, y_pred)
    fold accuracies.append(acc)
    print(f"Fold {fold} Accuracy: {acc:.4f}")

    fpr, tpr, _ = roc_curve(y_val, y_proba)
    auc = roc_auc_score(y_val, y_proba)
    plt.plot(fpr, tpr, label=f"Fold {fold} (AUC={auc:.2f})")

print("\nAverage CV Accuracy:", np.mean(fold accuracies))
# 4. Final Test Evaluation
y_pred_test = best_model.predict(X_test)
y_proba_test = best_model.predict_proba(X_test)[:, 1]

print("\n=== Performance Metrics on Test Set ===")
print("Accuracy :", accuracy_score(y_test, y_pred_test))
print("Precision:", precision_score(y_test, y_pred_test))
print("Recall   :", recall_score(y_test, y_pred_test))

print("F1 Score :", f1_score(y_test, y_pred_test))
print("ROC-AUC   :", roc_auc_score(y_test, y_proba_test))
print("\nClassification Report:\n", classification_report(y_test, y_pred_test))
# 5. Final ROC Curve for Test Set
fpr_test, tpr_test, _ = roc_curve(y_test, y_proba_test)
auc_test = roc_auc_score(y_test, y_proba_test)

plt.plot(fpr_test, tpr_test, color="black", linewidth=2.5, label=f"Test ROC (AUC={auc_test:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Gradient Boosting ROC Curve (5-Fold + Test)")
plt.legend()
plt.show()

```

5.XG Boosting

```
# 1. Load Dataset
data = load_breast_cancer()
X = data.data
y = data.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# 2. Hyperparameter Tuning
params = {
    "n_estimators": [50, 100, 150],
    "learning_rate": [0.01, 0.1, 0.2],
    "max_depth": [3, 5, 7]
}

grid = GridSearchCV(
    XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42),
    params,
    cv=5,
    scoring="accuracy",
    n_jobs=-1
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("Best Parameters:", grid.best_params_)
```

```
kf = KFold(n_splits=5, shuffle=True, random_state=42)
fold_accuracies = []
plt.figure(figsize=(8, 6))
for fold, (train_index, val_index) in enumerate(kf.split(X_train), start=1):
    X_tr, X_val = X_train[train_index], X_train[val_index]
    y_tr, y_val = y_train[train_index], y_train[val_index]

    best_model.fit(X_tr, y_tr)
    y_pred = best_model.predict(X_val)
    y_proba = best_model.predict_proba(X_val)[:, 1]
    acc = accuracy_score(y_val, y_pred)
    fold_accuracies.append(acc)
    print(f"Fold {fold} Accuracy: {acc:.4f}")

    fpr, tpr, _ = roc_curve(y_val, y_proba)
    auc = roc_auc_score(y_val, y_proba)
    plt.plot(fpr, tpr, label=f"Fold {fold} (AUC={auc:.2f})")

print("\nAverage CV Accuracy:", np.mean(fold_accuracies))

# 4. Final Test Evaluation
y_pred_test = best_model.predict(X_test)
y_proba_test = best_model.predict_proba(X_test)[:, 1]

print("\n=== Performance Metrics on Test Set ===")
print("Accuracy :", accuracy_score(y_test, y_pred_test))
print("Precision:", precision_score(y_test, y_pred_test))
print("Recall   :", recall_score(y_test, y_pred_test))
print("F1 Score  :", f1_score(y_test, y_pred_test))
```

```
# 5. Final ROC Curve for Test Set
fpr_test, tpr_test, _ = roc_curve(y_test, y_proba_test)
auc_test = roc_auc_score(y_test, y_proba_test)
plt.plot(fpr_test, tpr_test, color="black", linewidth=2.5, label=f"Test ROC (AUC={auc_test:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("XGBoost ROC Curve (5-Fold + Test)")
plt.legend()
plt.show()
```

6. Stacking Classifier

```
# 1. Load Dataset
data = load_breast_cancer()
X = data.data
y = data.target
# Split train-test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
# 2. Define Base Models and Stacking Classifier
base_estimators = [
    ('svm', SVC(probability=True, random_state=42)),
    ('nb', GaussianNB()),
    ('dt', DecisionTreeClassifier(random_state=42))
]
stack_model = StackingClassifier(
    estimators=base_estimators,
    final_estimator=LogisticRegression(max_iter=500, random_state=42),
    passthrough=False
)
# 3. Hyperparameter Tuning
params = {
    'svm__C': [0.1, 1, 10],
    'svm__kernel': ['linear', 'rbf'],
    'dt__max_depth': [3, 5, 7],
    'final_estimator__C': [0.1, 1, 10]
}

grid = GridSearchCV(
    stack_model,
    params,
    cv=5,
    scoring="accuracy",
    n_jobs=-1
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("Best Parameters:", grid.best_params_)
# 4. 5-Fold Cross-Validation + Fold-wise ROC
kf = KFold(n_splits=5, shuffle=True, random_state=42)
fold accuracies = []
plt.figure(figsize=(8, 6))
for fold, (train_index, val_index) in enumerate(kf.split(X_train), start=1):
    X_tr, X_val = X_train[train_index], X_train[val_index]
    y_tr, y_val = y_train[train_index], y_train[val_index]

    best_model.fit(X_tr, y_tr)
    y_pred = best_model.predict(X_val)
    y_proba = best_model.predict_proba(X_val)[: , 1]
    acc = accuracy_score(y_val, y_pred)
    fold accuracies.append(acc)
    print(f"Fold {fold} Accuracy: {acc:.4f}")

# ROC curve for each fold
```



```

fpr, tpr, _ = roc_curve(y_val, y_proba)
auc = roc_auc_score(y_val, y_proba)
plt.plot(fpr, tpr, label=f"Fold {fold} (AUC = {auc:.2f})")

print("\nAverage CV Accuracy:", np.mean(fold accuracies))
# 5. Final Evaluation on Test Set
y_pred_test = best_model.predict(X_test)
y_proba_test = best_model.predict_proba(X_test)[:, 1]

print("\n=== Performance Metrics on Test Set ===")
print("Accuracy :", accuracy_score(y_test, y_pred_test))
print("Precision:", precision_score(y_test, y_pred_test))
print("Recall   :", recall_score(y_test, y_pred_test))
print("F1 Score :", f1_score(y_test, y_pred_test))
print("ROC-AUC  :", roc_auc_score(y_test, y_proba_test))
print("\nClassification Report:\n", classification_report(y_test, y_pred_test))
# 6. Final ROC Curve for Test Set
fpr_test, tpr_test, _ = roc_curve(y_test, y_proba_test)
auc_test = roc_auc_score(y_test, y_proba_test)

plt.plot(fpr_test, tpr_test, color="black", linewidth=2.5, label=f"Test ROC (AUC = {auc_test:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Stacking Classifier ROC Curve (5-Fold + Test)")
plt.legend()
plt.show()

```

CONFUSION MATRIX AND ROC FOR EACH:

1.Decision Tree:

Evaluating Decision Tree Classifier (Criterion = gini)
=====

--- Testing max_depth = 2 ---

Accuracy : 0.8947

Precision: 0.9688

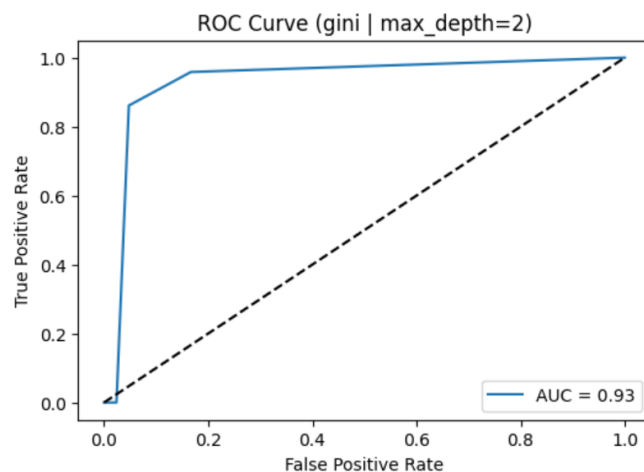
Recall : 0.8611

F1 Score : 0.9118

ROC-AUC : 0.9345

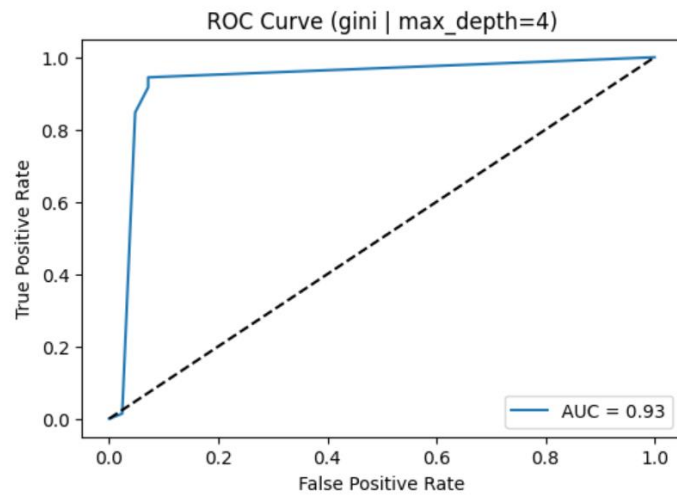
Classification Report:

	precision	recall	f1-score	support
0	0.80	0.95	0.87	42
1	0.97	0.86	0.91	72
accuracy			0.89	114
macro avg	0.88	0.91	0.89	114
weighted avg	0.91	0.89	0.90	114




```
--- Testing max_depth = 4 ---  
Accuracy : 0.9386  
Precision: 0.9577  
Recall   : 0.9444  
F1 Score : 0.9510  
ROC-AUC  : 0.9342
```

```
Classification Report:  
              precision    recall  f1-score   support  
  
     0       0.91         0.93         0.92         42  
     1       0.96         0.94         0.95         72  
  
 accuracy          0.94         0.94         0.94        114  
  macro avg       0.93         0.94         0.93        114  
 weighted avg     0.94         0.94         0.94        114
```



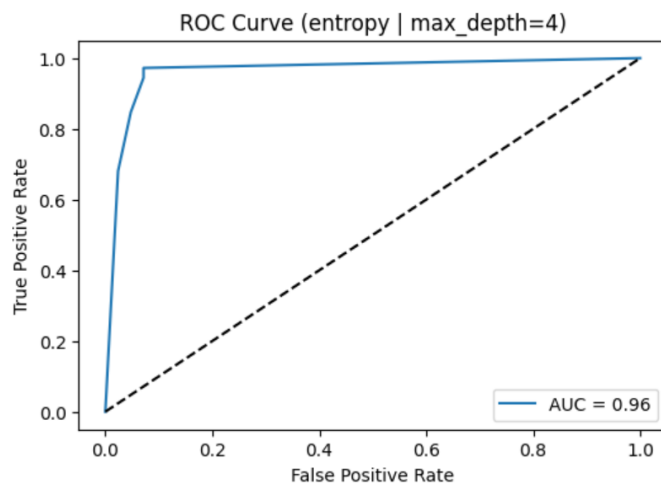
```
--- Testing max_depth = 2 ---
Accuracy : 0.8947
Precision: 0.9688
Recall   : 0.8611
F1 Score : 0.9118
ROC-AUC  : 0.9345
```

Classification Report:				
	precision	recall	f1-score	support
0	0.80	0.95	0.87	42
1	0.97	0.86	0.91	72
accuracy			0.89	114
macro avg	0.88	0.91	0.89	114
weighted avg	0.91	0.89	0.90	114



```
--- Testing max_depth = 4 ---
Accuracy : 0.9386
Precision: 0.9577
Recall   : 0.9444
F1 Score : 0.9510
ROC-AUC  : 0.9633
```

Classification Report:				
	precision	recall	f1-score	support
0	0.91	0.93	0.92	42
1	0.96	0.94	0.95	72
accuracy			0.94	114
macro avg	0.93	0.94	0.93	114
weighted avg	0.94	0.94	0.94	114



2.AdaBoost

```
Best Parameters: {'learning_rate': 1, 'n_estimators': 100}
Fold 1 Accuracy: 0.9560
Fold 2 Accuracy: 0.9560
Fold 3 Accuracy: 1.0000
Fold 4 Accuracy: 0.9890
Fold 5 Accuracy: 0.9341
```

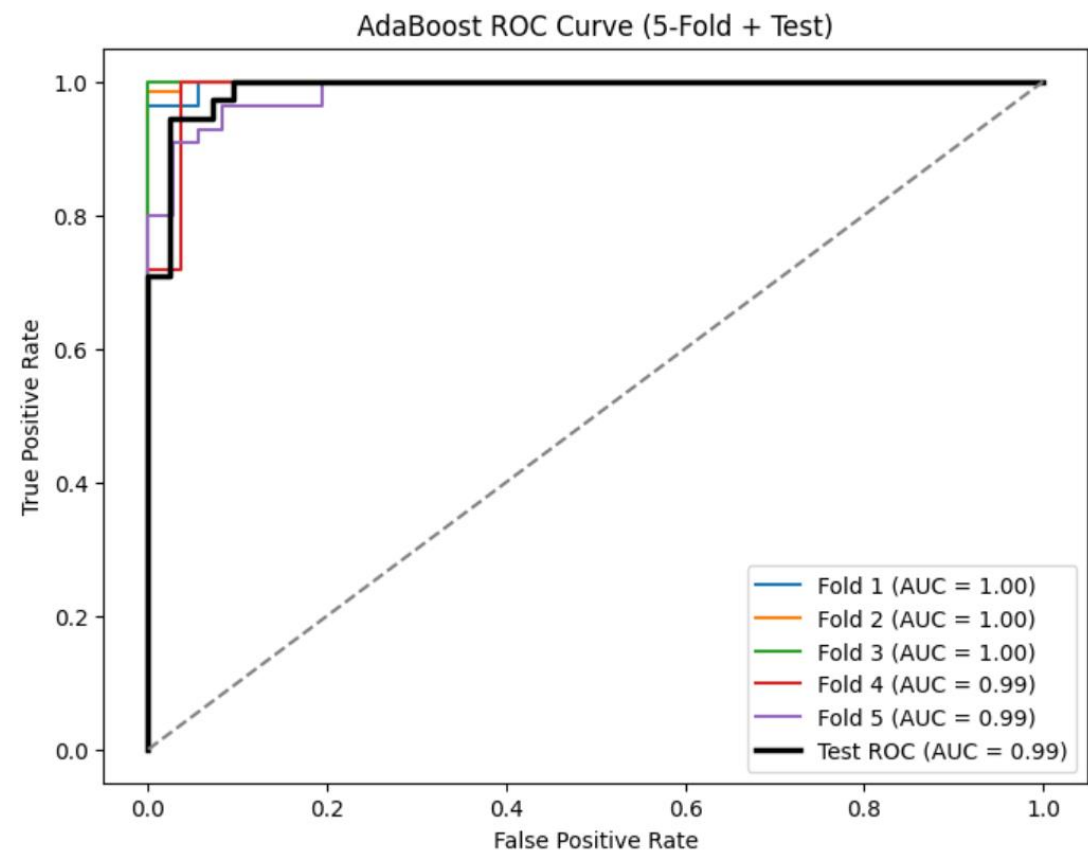

Average CV Accuracy: 0.9670329670329672

=== Performance Metrics on Test Set ===

```
Accuracy : 0.956140350877193
Precision: 0.9466666666666667
Recall    : 0.9861111111111112
F1 Score  : 0.9659863945578231
ROC-AUC   : 0.9897486772486772
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.90	0.94	42
1	0.95	0.99	0.97	72
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114



3. RandomForest

Best Parameters: {'max_depth': 7, 'min_samples_split': 5, 'n_estimators': 50}

Fold 1 Accuracy: 0.9341

Fold 2 Accuracy: 0.9670

Fold 3 Accuracy: 0.9670

Fold 4 Accuracy: 0.9670

Fold 5 Accuracy: 0.9341

Average CV Accuracy: 0.9538461538461538

=== Performance Metrics on Test Set ===

Accuracy : 0.9385964912280702

Precision: 0.9577464788732394

Recall : 0.9444444444444444


F1 Score : 0.951048951048951

ROC-AUC : 0.9914021164021164

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.93	0.92	42
1	0.96	0.94	0.95	72
accuracy			0.94	114
macro avg	0.93	0.94	0.93	114
weighted avg	0.94	0.94	0.94	114

Random Forest ROC Curve (5-Fold + Test)



4.GradientBoosting

Best Parameters: {'learning_rate': 0.2, 'max_depth': 3, 'n_estimators': 150}

Fold 1 Accuracy: 0.9011

Fold 2 Accuracy: 0.9780

Fold 3 Accuracy: 0.9670

Fold 4 Accuracy: 0.9890

Fold 5 Accuracy: 0.9341

Average CV Accuracy: 0.9538461538461538

=== Performance Metrics on Test Set ===

Accuracy : 0.9736842105263158

Precision: 0.96

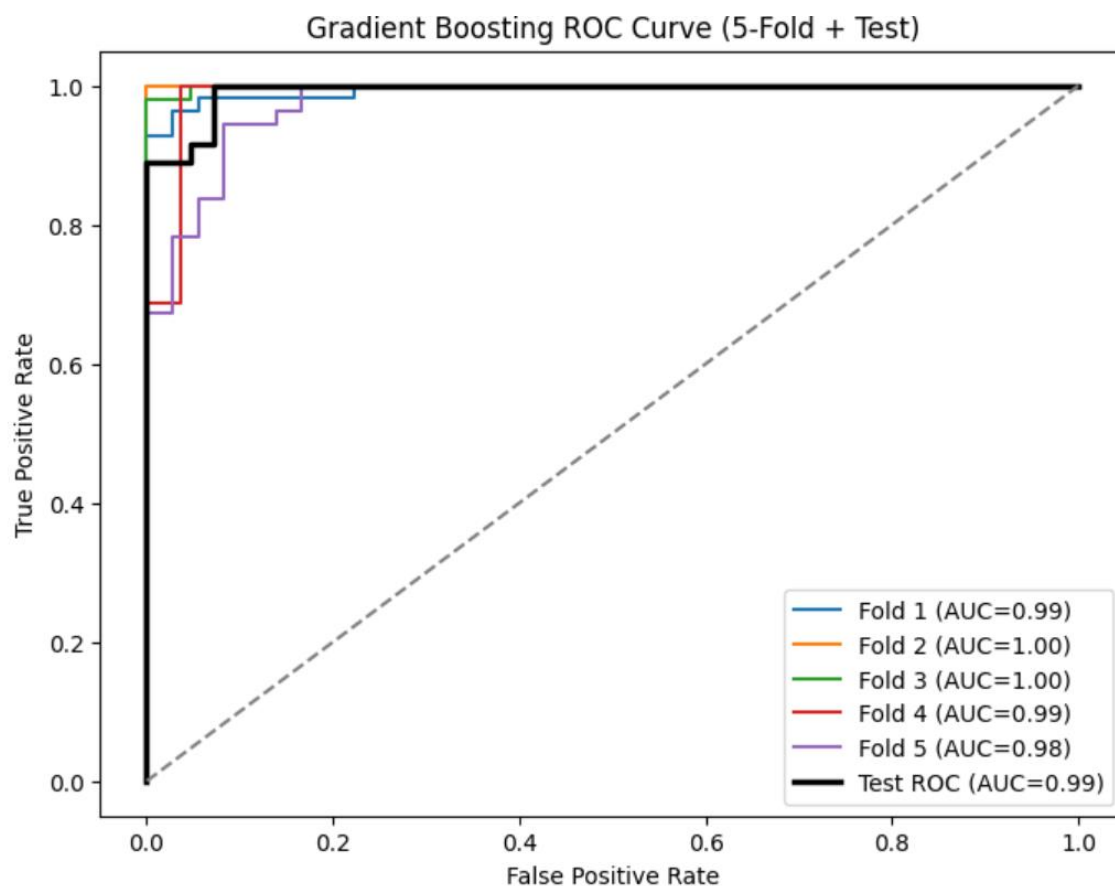
Recall : 1.0

F1 Score : 0.9795918367346939

ROC-AUC : 0.9927248677248678

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.93	0.96	42
1	0.96	1.00	0.98	72
accuracy			0.97	114
macro avg	0.98	0.96	0.97	114
weighted avg	0.97	0.97	0.97	114



5.XG Boosting

```
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [04:13:06] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
Best Parameters: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 150}
Fold 1 Accuracy: 0.9560
```

```
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [04:13:07] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [04:13:07] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
Fold 2 Accuracy: 0.9560
Fold 3 Accuracy: 0.9890
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [04:13:07] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183: UserWarning: [04:13:07] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
Fold 4 Accuracy: 0.9780
Fold 5 Accuracy: 0.9121
```

Average CV Accuracy: 0.9582417582417584

=== Performance Metrics on Test Set ===

Accuracy : 0.956140350877193

Precision: 0.9466666666666667

Recall : 0.9861111111111112


F1 Score : 0.9659863945578231

ROC-AUC : 0.9923941798941799

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.90	0.94	42
1	0.95	0.99	0.97	72
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

XGBoost ROC Curve (5-Fold + Test)



6.Stacking Model

```
Best Parameters: {'dt__max_depth': 5, 'final_estimator__C': 0.1, 'svm__C': 10, 'svm__kernel': 'linear'}
Fold 1 Accuracy: 0.9670
Fold 2 Accuracy: 0.9780
Fold 3 Accuracy: 0.9560
Fold 4 Accuracy: 0.9670
Fold 5 Accuracy: 0.9011
```

Average CV Accuracy: 0.9538461538461538

=== Performance Metrics on Test Set ===

Accuracy : 0.9473684210526315

Precision: 0.9459459459459459

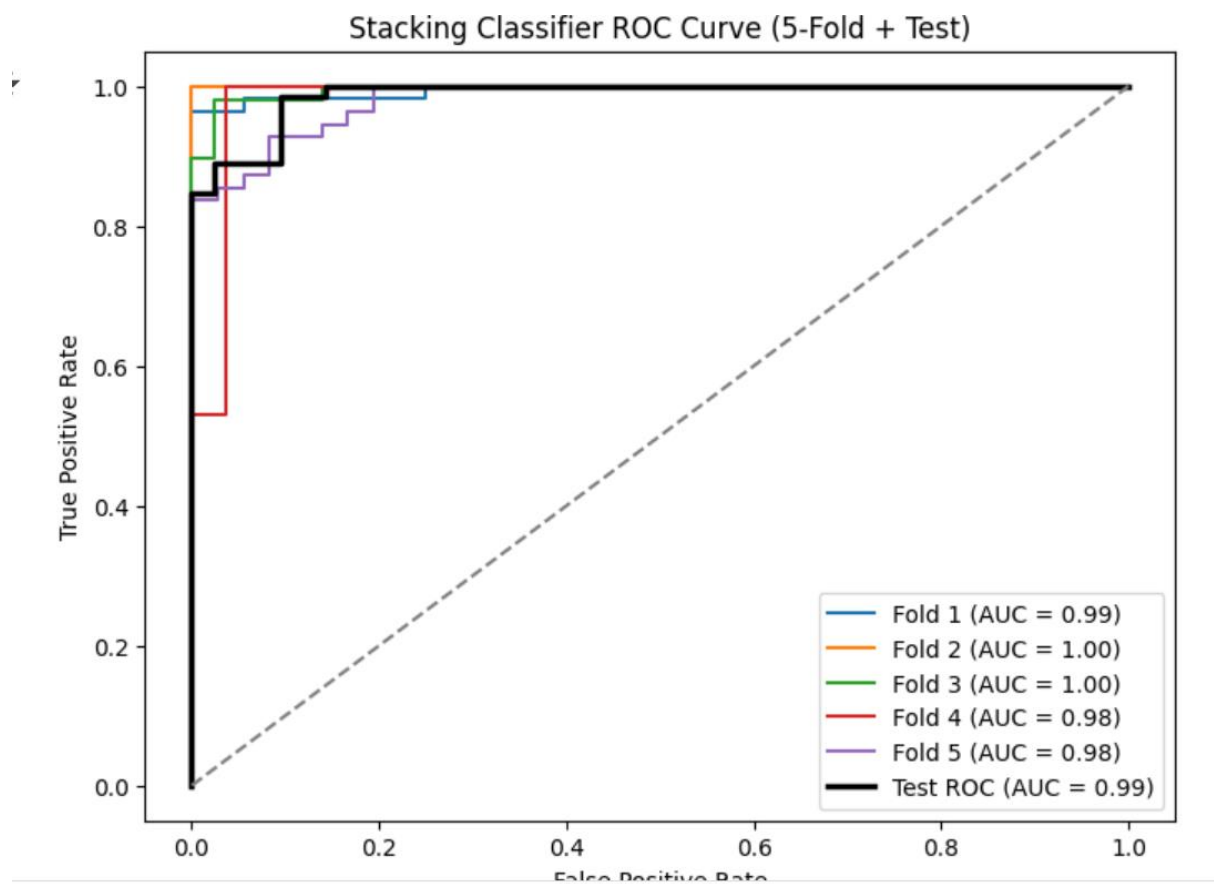
Recall : 0.9722222222222222

F1 Score : 0.958904109589041

ROC-AUC : 0.9877645502645503

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.90	0.93	42
1	0.95	0.97	0.96	72
accuracy			0.95	114
macro avg	0.95	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114



COMPARISON TABLES:

Decision Tree- Hyperparameter Tuning:

criterion	Max_depth	Accuracy	F1 Score
gini	2	0.8947	0.9118
gini	4	0.9386	0.9510
entropy	2	0.8947	0.9118
entropy	4	0.9386	0.9510

AdaBoost- Hyperparameter Tuning:

N_estimators	Learning_rate	Accuracy	F1 Score
100	1	0.967	0.965
150	1	0.967	0.973
150	0.15	0.958	0.966
350	0.15	0.956	0.965

GradientBoost- Hyperparameter Tuning:

N_estimators	Learning_rate	Max_depth	Accuracy	F1 Score
150	0.2	3	0.953	0.979
350	0.2	3	0.943	0.972
100	0.6	4	0.960	0.966

XGBoost- Hyperparameter Tuning:

N_estimators	Learning_rate	Max_depth	gamma	Accuracy	F1 Score
150	0.1	3	0	0.958	0.966
250	0.01	5	1	0.947	0.972
500	0.01	5	2	0.949	0.958

RandomForest- Hyperparameter Tuning:

N_estimators	Max_depth	Accuracy	F1 Score
50	7	0.9538	0.951
150	10	0.958	0.958
150	23	0.96	0.96

Stacked Ensemble - Hyperparameter Tuning:

Base Models	Final estimators	Accuracy	F1 Score
SVM, Naïve Bayes, Decision Tree	Logistic Regression	0.964	0.972
SVM, Naïve Bayes, Decision Tree	Random Forest	0.947	0.958
SVM, Decision Tree, KNN	Logistic Regression	0.952	0.971

Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Avg Accuracy
Decision Tree	0.9231	0.9780	1.0000	0.9780	0.9451	0.9648
AdaBoost	0.9341	0.9670	0.9670	0.9890	0.9341	0.9582
Gradient Boost	0.9451	0.9670	0.9780	0.9780	0.9341	0.9604
XGBoost	0.9341	0.9560	0.9670	0.9890	0.9011	0.9494
Random Forest	0.9670	0.9780	0.9560	0.9670	0.9011	0.9538

Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Avg Accuracy
Stacked Model	0.9341	0.9780	0.9780	0.9451	0.9341	0.9539

OBSERVATIONS:

1. Which model achieved the best validation accuracy among all six methods?

The AdaBoost classifier achieved the highest validation accuracy (96.7%), closely followed by Stacking (96.4%) and Random Forest (96.0%).

2. How does Decision Tree performance compare to ensemble methods?

Ensemble models significantly outperformed Decision Trees. While the Decision Tree peaked at 93.8%, ensemble techniques achieved 95–96.7%, showing better stability and predictive power.

3. Did the Random Forest benefit from tuning max_depth or n_estimators?

Yes, tuning n_estimators and max_depth improved Random Forest's performance. Optimal results were obtained with 150 trees and max_depth=23, achieving 96.0% accuracy.

4. Which model showed the best generalization? Any overfitting?

Stacking and AdaBoost showed the best generalization. A standalone Decision Tree slightly overfitted, while ensemble methods like Random Forest and XGBoost handled overfitting well.

5. Did stacking improve performance over base models?

Stacking improved performance compared to individual base models. Using Logistic Regression as the final estimator in stacking achieved 96.4%, which outperformed SVM, KNN, and Decision Tree alone

CONCLUSIONS:

1. We have learnt that AdaBoost consistently outperformed other classifiers, achieving the highest accuracy of 96.7%, while Stacking also performed very well with 96.4%, showing the strength of ensemble learning.
2. We have learnt that hyperparameter tuning plays a crucial role in improving performance — for example, increasing `n_estimators` and adjusting `learning_rate` boosted AdaBoost, while tuning `max_depth` and `n_estimators` enhanced Random Forest accuracy.
3. We have learnt that different algorithms provide trade-offs between accuracy, stability, and generalization — standalone Decision Trees tend to overfit, whereas ensemble methods like GradientBoost, XGBoost, and Stacking achieve better generalization and higher overall performance.